hypercluster ? end-to-end Kubernetes cluster lifecycle wrapper for Kubespray.

Usage: hypercluster [OPTIONS] <command> [ARGS]

Commands:

```
 init             Clone Kubespray and generate inventory + group_vars
                Use init --wizard to create cluster.conf interactively first
                Use init --profile <name> to copy config/profiles/<name>.conf.example
 wizard           Interactive cluster.conf generator (no existing config required)
 preflight        Verify local tools and SSH connectivity to all nodes (--json)
 deploy           Full cluster deployment (preflight + init + cluster.yml)
 scale            Add worker nodes listed in cluster.conf (scale.yml)
 upgrade          Upgrade Kubernetes to the version in cluster.conf
 remove-node      Drain and remove a single node
 fetch-kubeconfig  Copy admin.conf from the first control-plane node
 merge-kubeconfig  Merge the cluster kubeconfig into ~/.kube/config
 install-kubevirt  Install KubeVirt operator + CDI (also runs after deploy if KUBEVIRT_ENABLED=true)
 test             Run platform smoke tests (k3s path; optional --full, --json)
 destroy          Tear down the cluster (reset.yml or k3s uninstall)
 vms              List KubeVirt VMs + template catalog (--json)
 vm create        Validate template and preview VM create (--template, --name, --json)
 workloads        List pods and services (--json)
 discover         Discover VMs/hosts (libvirt, Proxmox, VMware, ssh-scan, csv) (--json)
 status           Node status (optional --json for machine-readable output)
 health           Connect + report nodes, pods, services, VMs (--json; --with-test for k3s smoke)
 unlock           Remove stale per-cluster lock after a crash
 lab setup        Secondary disk + memory tune (from cluster.conf LAB_* vars)
 forge-handoff    Emit JSON handoff for Forge Day-2 import (--json)
 lab storage      Mount LAB_SECONDARY_DISK, patch local-path, StorageClass
 lab tune         Scale down HA replicas on single-node labs (--restore to undo)
```

Options:

```
 -c <file>   Path to cluster config (default: ./cluster.conf)
 -f          Force overwrite of existing inventory
 -n          Dry run ? print what would be executed, skip actual changes
 -v          Verbose Ansible output
 -V          Show version
 -h          Show this help
?? defaults ????????????????????????????????????????????????????????????
?? parse options ???????????????????????????????????????????????????????
?? load libraries ??????????????????????????????????????????????????????
shellcheck source=lib/log.sh
shellcheck source=lib/wizard.sh
shellcheck source=lib/progress.sh
shellcheck source=lib/postdeploy.sh
shellcheck source=lib/preflight.sh
shellcheck source=lib/inventory.sh
shellcheck source=lib/kubespray.sh
```

shellcheck source=lib/kubeconfig.sh
shellcheck source=lib/k3s.sh
shellcheck source=lib/platform-stack.sh
shellcheck source=lib/platform-tools.sh
shellcheck source=lib/platform-test.sh
shellcheck source=lib/discover.sh
shellcheck source=lib/rke2.sh
shellcheck source=lib/microk8s.sh
shellcheck source=lib/talos.sh
shellcheck source=lib/vms.sh
shellcheck source=lib/workloads.sh
shellcheck source=lib/health.sh
shellcheck source=lib/forge-handoff.sh
?? discover (no cluster.conf required) ??????????????????????????????????????
?? wizard (no cluster.conf required) ??
?? load cluster config ??
shellcheck source=/dev/null
Validate required config variables
kubeadm maps to Kubespray path
?? lock helpers ???
unlock skips lock acquisition
Resolve KUBESPRAY_DIR: absolute paths pass through; relative paths anchor to repo or app data dir.
shellcheck source=lib/runtime-paths.sh
?? commands ??
?? dispatch ??